

### Practical 1:

1. Import all the required Python Libraries.
2. Locate an open source data from the web (e.g., <https://www.kaggle.com>). Provide a clear description of the data and its source (i.e., URL of the web site).
3. Load the Dataset into pandas dataframe.
4. Data Preprocessing: check for missing values in the data using pandas `isnull()`, `describe()` function to get some initial statistics.  
Provide variable descriptions. Types of variables etc. Check the dimensions of the data frame.
5. Data Formatting and Data Normalization: Summarize the types of variables by checking the data types (i.e., character, numeric, integer, factor, and logical) of the variables in the data set. If variables are not in the correct data type, apply proper type conversions.
6. Turn categorical variables into quantitative variables in Python

```
import pandas as pd
import numpy as np
```

```
df = pd.read_csv('Pr1.csv')
df
```

```
df.describe()
```

```
df.isnull()
```

```
df.isnull().sum()
```

```
df.info()
```

```
df.shape
```

```
df.dtypes
```

```
df = df.replace(to_replace = np.nan, value = 0.0)
df
```

```
df['Gender'].replace(to_replace = ['F', 'M'], value = [0, 1], inplace = True)
df
```

```
df2 = df.groupby('Division')
df2.get_group('A')
```

## Practical 2:

Create an “Academic performance” dataset of students and perform the following operations using Python.

1. Scan all variables for missing values and inconsistencies. If there are missing values and/or inconsistencies, use any of the suitable techniques to deal with them.
2. Scan all numeric variables for outliers. If there are outliers, use any of the suitable techniques to deal with them.
3. Apply data transformations on at least one of the variables. The purpose of this transformation should be one of the following reasons:  
to change the scale for better understanding of the variable, to convert a non-linear relation into a linear one, or to decrease the skewness and convert the distribution into a normal distribution.

Reason and document your approach properly.

```
import numpy as np
import pandas as pd

df = pd.read_csv("Pr2.csv")
df

df.describe()

df.info()

df.isnull().sum()

avg_val = df["Discussion"].astype("float").mean()
avg_val

df['Discussion'].replace(to_replace = np.nan, value = avg_val, inplace = True)
df.isnull().sum()

import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats

sns.boxplot(x=df['AnnouncementsView'])
plt.show()

sns.regplot(x='Sno', y='AnnouncementsView', data=df)
plt.show()

z = np.abs(stats.zscore(df['AnnouncementsView']))
z

threshold = 3
print(np.where(z > threshold))

z[419]

avg_val = int(df[df.index != 419]["AnnouncementsView"].mean())
```

```
avg_val
```

```
df.at[419, 'AnnouncementsView'] = avg_val  
sns.boxplot(x=df['AnnouncementsView'])  
plt.show()
```

```
from sklearn.preprocessing import MinMaxScaler  
scaler = MinMaxScaler()  
discussion_values = df['Discussion'].values.reshape(-1, 1)  
df['Discussion_scaled'] = scaler.fit_transform(discussion_values)
```

```
plt.figure(figsize=(12, 6))
```

```
plt.subplot(1, 2, 1)  
sns.histplot(df['Discussion'], kde=True, color='blue')  
plt.title('Original Distribution')
```

```
plt.subplot(1, 2, 2)  
sns.histplot(df['Discussion_scaled'], kde=True, color='green')  
plt.title('MinMaxScaled Distribution')
```

```
plt.tight_layout()  
plt.show()
```

### Practical 3:

#### Descriptive Statistics - Measures of Central Tendency and variability

Perform the following operations on any open source dataset (e.g., data.csv)

1. Provide summary statistics (mean, median, minimum, maximum, standard deviation) for a dataset (age, income etc.) with numeric variables grouped by one of the qualitative (categorical) variable. For example, if your categorical variable is age groups and quantitative variable is income, then provide summary statistics of income grouped by the age groups. Create a list that contains a numeric value for each response to the categorical variable.

2. Write a Python program to display some basic statistical details like percentile, mean, standard deviation etc. of the species of

'Iris-setosa', 'Iris-versicolor' and 'Iris-versicolor' of iris.csv dataset.

Provide the codes with outputs and explain everything that you do in this step.

```
import pandas as pd
import numpy as np
import statistics as st
```

```
df = pd.read_csv("Mall_Customers.csv")
df
```

```
df.mean(numeric_only = True)
```

```
#Mean of a specific column:
df.loc[:, 'Age'].mean()
```

```
#Row-wise mean:
df.mean(numeric_only = True, axis=1)
```

```
df.median(numeric_only = True)
```

```
df.mode()
```

```
df.min()
```

```
df.max()
```

```
df.std(numeric_only = True)
```

```
df.groupby(['Genre'])['Age'].mean()
```

```
df.rename(columns= {'Annual Income (k$)': 'Income'}, inplace= True)
df
```

```
df.groupby(['Genre']).Income.mean()
```

```
df_iris = pd.read_csv("Iris.csv")
df_iris
```

```
print('Iris-setosa')
```

```
setosa = df_iris['Species'] == 'Iris-setosa'
print(df_iris[setosa].describe())
print("\nIris-versicolor")
versicolor = df_iris['Species'] == 'Iris-versicolor'
print(df_iris[versicolor].describe())
print("\nIris-virginica")
virginica = df_iris['Species'] == 'Iris-virginica'
print(df_iris[virginica].describe())
```

## Practical 4:

### Data Analytics I

Create a Linear Regression Model using Python/R to predict home prices using Boston Housing Dataset (<https://www.kaggle.com/c/boston-housing>). The Boston Housing dataset contains information about various houses in Boston through different parameters. There are 506 samples and 14 feature variables in this dataset.

The objective is to predict the value of prices of the house using the given features.

```
#load the Boston Housing DataSet
dataset <- read.csv("Boston.csv")

data <- head(dataset,100)
# Extract the columns for RM and MEDV

x <- data$rm
y <- data$medv

# Fit a linear regression model
fit <- lm(y ~ x)

#plot the data and the regression line
plot(x,y,main="Linear Regression Model for Boston Housing Dataset",
      xlab = "Average Number of Rooms per Dwelling",
      ylab = "Median Value of Owner-Occupied Homes",
      pch = 16, col = "blue")
abline(fit,col = "red")
```

## Practical 5:

### Data Analytics II

1. Implement logistic regression using Python/R to perform classification on Social\_Network\_Ads.csv dataset.
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset

```
#load the dataset
```

```
data <- read.csv("Social Network Ads.csv")
```

```
#Extract the columns
```

```
x <- data$EstimatedSalary
```

```
y <- data$Purchased
```

```
#Fit a regression model
```

```
fit <- glm(y ~ x, data = data, family="binomial")
```

```
#Create a graph for regression
```

```
plot(x,y,main = "Logistic Regression", xlab = "Estimated Salary", ylab = "Ads Purchased",  
      col = ifelse(y == 1,"red","blue"))
```

```
abline(model, col = "red")
```

```
curve(predict(fit, data.frame(x), type = "response"), add = TRUE, col = "black", lwd = 2)  
dev.off()
```

```
# Predict the probabilities
```

```
predicted_prob <- predict(fit, type = "response")
```

```
# Convert probabilities to binary predictions
```

```
predicted <- ifelse(predicted_prob >= 0.5, 1, 0)
```

```
# Create a confusion matrix
```

```
conf_matrix <- table(Actual = data$Purchased, Predicted = predicted)
```

```
# Compute evaluation metrics
```

```
TP <- conf_matrix[1, 1]
```

```
FN <- conf_matrix[1, 2]
```

```
FP <- conf_matrix[2, 1]
```

```
TN <- conf_matrix[2, 2]
```

```
Accuracy <- sum(diag(conf_matrix)) / sum(conf_matrix)
```

```
Error_rate <- 1 - Accuracy
```

```
Precision <- TP / (TP + FP)
```

```
Recall <- TP / (TP + FN)
```

```
# Print confusion matrix and evaluation metrics
```

```
print("Confusion Matrix:")
```

```
print(conf_matrix)
```

```
print(c("Accuracy:", Accuracy))
```

```
print(c("Error rate:", Error_rate))
```

```
print(c("Precision:", Precision))
```

```
print(c("Recall:", Recall))
```

## Practical 6:

### Data Analytics III

1. Implement Simple Naïve Bayes classification algorithm using Python/R on iris.csv dataset.
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset.

```
library(e1071)
```

```
iris <- read.csv("Iris.csv")
```

```
trainIndex <- sample(1:nrow(iris), 0.7*nrow(iris), replace = FALSE)
```

```
train <- iris[trainIndex,]
```

```
test <- iris[-trainIndex,]
```

```
# Creating the Naive Bayes model
```

```
model <- naiveBayes(Species ~ ., data = train)
```

```
# Making predictions on the test set
```

```
predictions <- predict(model, test)
```

```
print("Predictions")
```

```
print(predictions)
```

```
# Creating the confusion matrix
```

```
conf_matrix <- table(test$Species, predictions)
```

```
print("Confusion Matrix")
```

```
print(conf_matrix)
```

```
# Calculating the evaluation metrics
```

```
TP <- conf_matrix[1,1]
```

```
FP <- conf_matrix[2,1] + conf_matrix[3,1]
```

```
TN <- conf_matrix[2,2] + conf_matrix[3,2] + conf_matrix[1,3] + conf_matrix[2,3] +  
conf_matrix[3,3]
```

```
FN <- conf_matrix[1,2] + conf_matrix[1,3]
```

```
Accuracy <- sum(diag(conf_matrix))/sum(conf_matrix)
```

```
Error_rate <- 1 - Accuracy
```

```
Precision <- TP/(TP+FP)
```

```
Recall <- TP/(TP+FN)
```

```
print(c("Accuracy: ",Accuracy))
```

```
print(c("Error_rate: ",Error_rate))
```

```
print(c("Precision: ",Precision))
```

```
print(c("Recall: ",Recall))
```



## Practical 7:

### Text Analytics

1. Extract Sample document and apply following document preprocessing methods: Tokenization, POS Tagging, stop words removal, Stemming and Lemmatization.
2. Create representation of document by calculating Term Frequency and Inverse Document Frequency.

```
sentence1 = "I will walk 500 miles and I would walk 500 more. Just to be the man who walks a thousand miles to fall down at your door!"
```

```
sentence2 = "I played the play playfully as the players were playing in the play with playfulness"
```

```
from nltk import word_tokenize
print("Tokenized words:", word_tokenize(sentence1))
```

```
from nltk import pos_tag
token = word_tokenize(sentence1) + word_tokenize(sentence2)
tagged = pos_tag(token)
print("Tagging Parts of Speech:", tagged)
```

```
from nltk.corpus import stopwords
stop_words = stopwords.words('english')
token = word_tokenize(sentence1)
cleaned_token = []
for word in token:
    if word not in stop_words:
        cleaned_token.append(word)
print('Unclean version:', token)
print('\nCleaned version:', cleaned_token)
```

```
from nltk.stem import PorterStemmer
stemmer = PorterStemmer()
stemmed_tokens = [stemmer.stem(token) for token in cleaned_token]
print(stemmed_tokens)
```

```
from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
lemmatized_tokens = [lemmatizer.lemmatize(token) for token in cleaned_token]
print(lemmatized_tokens)
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
tfidf_vectorizer = TfidfVectorizer()
tfidf_matrix = tfidf_vectorizer.fit_transform([sentence1])
print(tfidf_matrix.toarray())
print("\n# FEATURED NAMES \n")
print(tfidf_vectorizer.get_feature_names_out())
```

## Practical 8:

### Data Visualization I

1. Use the inbuilt dataset 'titanic'. The dataset contains 891 rows and contains information about the passengers who boarded the unfortunate Titanic ship. Use the Seaborn library to see if we can find any patterns in the data.
2. Write a code to check how the price of the ticket (column name: 'fare') for each passenger is distributed by plotting a histogram.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

data = pd.read_csv('Titanic.csv')
data

data.describe()

data.isnull().sum()

data['Age'] = data['Age'].fillna(data['Age'].mean(numeric_only = True))
data['Cabin'] = data['Cabin'].fillna(data['Cabin'].mode()[0])
data['Embarked'] = data['Embarked'].fillna(data['Embarked'].mode()[0])

data.isnull().sum()

sns.countplot(x='Survived',data=data)

sns.countplot(x='Pclass',data=data)

sns.countplot(x='Embarked',data=data)

sns.countplot(x='Sex',data=data)

sns.boxplot(data['Age'])

sns.boxplot(data['Fare'])

sns.boxplot(data['Pclass'])

sns.boxplot(data['SibSp'])

sns.catplot(x= 'Pclass', y = 'Age', data=data, kind = 'box')

sns.catplot(x= 'Pclass', y = 'Fare', data=data, kind = 'strip')

sns.catplot(x= 'Sex', y = 'Fare', data=data, kind = 'strip')

sns.catplot(x= 'Sex', y = 'Age', data=data, kind = 'strip')

sns.pairplot(data)
```

```
sns.scatterplot(x = 'Fare', y = 'Pclass', hue = 'Survived', data = data)
```

```
sns.scatterplot(x = 'Survived', y = 'Fare', data = data)
```

```
sns.distplot(data['Age'])
```

```
sns.distplot(data['Fare'])
```

```
sns.jointplot(x = "Survived", y = "Fare", kind = "scatter", data = data)
```

```
sns.heatmap(data.corr(numeric_only = True), cmap="YlGnBu")  
plt.title('Correlation')
```

```
sns.catplot(x='Pclass', y='Fare', data=data, kind='bar')
```

```
plt.hist(data['Fare'])
```

## Practical 9:

### Data Visualization II

1. Use the inbuilt dataset 'titanic' as used in the above problem. Plot a box plot for distribution of age with respect to each gender along with the information about whether they survived or not. (Column names : 'sex' and 'age')
2. Write observations on the inference from the above statistics.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

data = pd.read_csv('Titanic.csv')
data

data.describe()

data.isnull().sum()

data['Age'] = data['Age'].fillna(data['Age'].mean(numeric_only = True))
data['Cabin'] = data['Cabin'].fillna(data['Cabin'].mode()[0])
data['Embarked'] = data['Embarked'].fillna(data['Embarked'].mode()[0])
data.isnull().sum()

sns.boxplot(x='Sex', y='Age', data=data, hue="Survived")
```

## Practical 10:

### Data Visualization III

Download the Iris flower dataset or any other dataset into a DataFrame. Scan the dataset and give the inference as:

1. List down the features and their types (e.g., numeric, nominal) available in the dataset.
2. Create a histogram for each feature in the dataset to illustrate the feature distributions.
3. Create a boxplot for each feature in the dataset.
4. Compare distributions and identify outliers.

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
df = pd.read_csv('Iris.csv')
df
```

```
df.isnull().sum()
```

```
df.info()
```

```
np.unique(df["Species"])
```

```
df.describe()
```

```
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 8))
```

```
sns.histplot(df['SepalLengthCm'], kde=True, color='blue', ax = axes[0,0])
axes[0,0].set_title('Sepal Length')
```

```
sns.histplot(df['SepalWidthCm'], kde=True, color='green', ax = axes[0,1])
axes[0,1].set_title('Sepal Width')
```

```
sns.histplot(df['PetalLengthCm'], kde=True, color='yellow', ax = axes[1,0])
axes[1,0].set_title('Petal Length')
```

```
sns.histplot(df['PetalWidthCm'], kde=True, color='red', ax = axes[1,1])
axes[1,1].set_title('Petal Width')
```

```
plt.tight_layout()
plt.show()
```

```
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 8))
```

```
sns.boxplot(df['SepalLengthCm'], color='blue', ax = axes[0,0])
axes[0,0].set_title('Sepal Length')
```

```
sns.boxplot(df['SepalWidthCm'], color='green', ax = axes[0,1])
axes[0,1].set_title('Sepal Width')
```

```
sns.boxplot(df['PetalLengthCm'], color='yellow', ax = axes[1,0])
axes[1,0].set_title('Petal Length')
```

```
sns.boxplot(df['PetalWidthCm'], color='red', ax = axes[1,1])  
axes[1,1].set_title('Petal Width')
```

```
plt.tight_layout()  
plt.show()
```

```
data_to_plot = [df[x] for x in df.columns[1:-1]]  
fig, axes = plt.subplots(1, figsize=(12,8))  
bp = axes.boxplot(data_to_plot)
```

Practical 13:  
SCALA Program for word count:

```
val text=sc.textFile("my_text.txt")
```

```
val counts=text.flatMap(line=>line.split(" ")).map(word=>(word,1)).reduceByKey(_+_)
```

```
counts.collect
```